

Programovanie v operačných systémoch


02 - Filesystem IO

Jozef Šiška



Department of Applied Informatics
Comenius University in Bratislava

2025/2026



- ▶ ANSI C, ISO C (C89, C99, C11)
- ▶ POSIX ([IEEE](#))
- ▶ [Single Unix Specification](#) (The Open Group Base Specifications [2018](#) identical to POSIX, other "standards" for what a UNIX system should be / should have)
- ▶ Linux Standard Base (LSB) - linux specific, for binary interoperability between linux distros (filesystem hierarchy, standard libraries, file locations)

- ▶ VFS (virtual filesystem)
- ▶ mounted "real" filesystems (`mount(2)`)
- ▶ files
 - ▶ name, data, metadata
 - ▶ inode
 - ▶ `open` (`creat`), `close`, `read`, `write`, `stat`, ...
- ▶ directories ("folders")
 - ▶ list of entries (files, directories)
 - ▶ POSIX C: `opendir(3)`, `readdir(3)`, `closedir(3)`
 - ▶ syscalls: (`open(2)`, `close(2)`,) `getdents(2)`, `readdir(2)`, no C wrapper

Digression: man

man man

- ▶ 1 User Commands
- ▶ 2 System Calls
- ▶ 3 C Library Functions
- ▶ ...

```
man open      # open(2)
man close     # close(2)
man 1 mkdir   # mkdir(1): command
man 2 mkdir   # mkdir(2): C function / syscall
man 1 printf  # printf(1): command
man 3 printf  # printf(2): C function
```

Maybe not a file

- ▶ Regular files, directories
- ▶ Block and character devices
- ▶ Symbolic links
- ▶ Pipes
- ▶ Sockets

<code>mknod(2)</code>	Create "special" (device) files
<code>symlink(2)</code>	Create a symbolic link
<code>pipe(2)</code> , <code>pipe2(2)</code>	Creates a pipe
<code>socket(2)</code>	Create a network socket

Filesystem related calls

<code>fd = open(path, flags)</code>	Open/create a file
<code>close(fd)</code>	Close a file
<code>count = read(fd, &buf, count)</code>	Read from file
<code>count = write(fd, &buf, count)</code>	Write to a file
<code>pos = lseek(fd, offset, whence)</code>	Change position
<code>stat(path, &buf), fstat(fd, &buf)</code>	Get file (status) information
<code>access(name, mode)</code>	Check accessibility / existence
<code>chmod(n, mode), chown(n, u, g)</code>	Change permissions, owner
<code>utime(name, times)</code>	Change file access / modification times
<code>umask(mode)</code>	Change file creation mode mask
<code>rename(old, new)</code>	Rename file
<code>mkdir(name, mode), rmdir(name)</code>	Create / remove empty directory
<code>link(n1, n2), unlink(name)</code>	Create link, remove dir entry
<code>mount(special, name, flags)</code>	Mount fs
<code>umount(special)</code>	Unmount fs
<code>sync(), syncfs(), fsync(), fdatasync()</code>	Synchronize file(system) to disk
<code>chdir(dirname), chroot(dirname)</code>	Change current / root directory

... and more

`dup(fd)`, `dup2(fd,newfd)`, `dup3(..., flags)`

Duplicate fd

`fcntl(fd, cmd, args...)`

File "operations"

`ioctl(fd, request, args...)`

"Special" operations on fd

... and more

dup(fd), dup2(fd,newfd), dup3(..., flags)	Duplicate fd
fcntl(fd, cmd, args...)	File "operations"
ioctl(fd, request, args...)	"Special" operations on fd

```
int fdIn = open("a.txt", O_RDONLY);
if (fdIn == -1) {
    perror("open"); // prints an error depending on errno
    exit(EXIT_FAILURE);
}

int fdOut;
if ((fdOut = open("a.txt", O_WRONLY | O_CREAT)) == -1) {
    ...
}
```

Aside: error handling

This class will also be about religiously handling every error / return value.

However: because space is precious on slides, examples on slides will almost always ommit it ; -).

Get output of a program

```
int pipefd[2];
pid_t pid;

pipe(pipefd);

pid = fork();
if (pid == 0) { // child
    dup2(pipefd[1], 1);
    execvp(argv[1], argv+1);
} else { // parent
    char buffer[1024];
    ssize_t br = read(pipefd[0], buffer, sizeof(buffer));
    printf("Output: \%.s\n", (int)br, buffer);
    int wstatus;
    if (waitpid(pid, &wstatus, 0) == -1) {
        perror("waitpid");
        exit(EXIT_FAILURE);
    }
    if (WIFEXITED(wstatus)) {
        printf("Exited: \%.d\n", WEXITSTATUS(wstatus));
    }
}

return 0;
```

Non-blocking IO

- ▶ `open(..., O_NONBLOCK)`

- ▶ `fcntl` on already open fds

```
int flags = fcntl(fd, F_GETFL, 0); // check for -1
fcntl(fd, F_SETFL, flags | O_NONBLOCK); // check for -1
```

- ▶ `read/write`: return `EAGAIN`, `EWOULDBLOCK`
- ▶ `select`, `poll`, `epoll`: Wait for events on file descriptors
- ▶ `ioctl(fd, FIONREAD, &bytes_available)`

Filesystem io always returns the full file size as available, and will block (if data needs to be read from disk etc.).