

# Programovanie v operačných systémoch

## 05 - Network

Jozef Šiška



Department of Applied Informatics  
Comenius University in Bratislava

2025/2026



# Networking Recap

- ▶ HW
- ▶ Ethernet
- ▶ IP
  - ▶ TCP
  - ▶ UDP
  - ▶ DNS
- ▶ Local sockets (Unix sockets)

# Socket

Create a socket (not yet connected to anything)

```
int socket(int domain, int type, int protocol);
```

## Domain

AF_UNIX, AF_LOCAL	local sockets (man 7 unix)
AF_INET	IPv4 (man 7 ip)
AF_INET6	IPv6 (man 7 ipv6)
AF_NETLINK	kernel userspace interface (man 7 netlink)
AF_IPX, AF_X25,...	

## Type

SOCK_STREAM	reliable byte stream (i.e. TCP)
SOCK_DGRAM	connectionless, unreliable messages (UDP)
SOCK_SEQPACKET, SOCK_RAW, ...	

`man 2 socket`

`man 7 {unix,ip,ipv6}`



## Server

- ▶ Bind a socket to an address

```
int bind(int sockfd, const struct sockaddr *addr,  
         socklen_t addrlen);
```

- ▶ Listen on the socket for incoming connections

```
int listen(int sockfd, int backlog);
```

- ▶ Accept a connection

```
int accept(int sockfs, struct sockaddr *addr,  
          socklen_t *addrlen);
```

## Client

```
int connect(int sockfd, const struct sockaddr *addr,  
            socklen_t addrlen);
```

```
man 2 {bind,listen,accept,connect}
```

# Reading, writing, closing

Read	Write	
read	write	plain read/write
recv	send	specify additional flags
recvfrom	sendto	get / specify peer address (i.e. UDP packets)
recvmsg	sendmsg	readv/writev style, additional data
shutdown		close (one direction of) a connection
close		close (dispose of) the socket

# Addresses

A general "some address" type (man 2 bind):

```
struct sockaddr {  
    sa_family_t sa_family;  
    char        sa_data[14];  
};
```

IPv4 address (man 7 ip, IPv6 is similar):

```
struct sockaddr_in {  
    sa_family_t    sin_family; /* address family: AF_INET */  
    in_port_t      sin_port;   /* port in network byte order */  
    struct in_addr sin_addr;  /* internet address */  
};  
  
/* Internet address. */  
struct in_addr {  
    uint32_t        s_addr;    /* address in network byte order */  
};
```

Need to cast between types:

```
struct sockaddr_in  addr;  
/* set the fields, open socket */  
ret = bind(sockfd, (struct sockaddr *) &addr, sizeof(addr));
```

# Obtaining, printing addresses

## Any address (for server)

```
struct sockaddr_in addr;  
addr.sin_family = AF_INET;  
addr.sin_addr.s_addr = INADDR_ANY;
```

## Network vs host order (ports)

```
uint16_t portno = 1234;  
addr.sin_port = htons(portno);  
portno = ntohs(addr.sin_port);
```

## Convert IPv4 address to sockaddr\_in

```
ret = inet_aton("127.0.0.1", &addr.sin_addr);  
ret = inet_pton(AF_INET, "127.0.0.1", &addr.sin_addr);  
ret = inet_pton(AF_INET6, "::1", &addr.sin6_addr);
```

## Convert sockaddr\_in to IPv4 address

```
printf("%s\n", inet_ntoa(addr.sin_addr));  
char str[INET_ADDRSTRLEN];  
ret = inet_ntop(AF_INET, &addr.sin_addr, str, len);
```

# Resolving DNS addresses

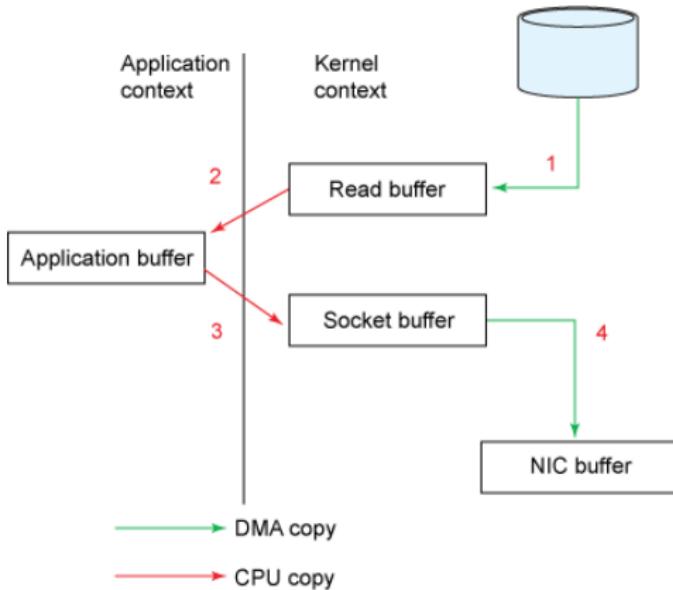
## gethostbyname

```
struct sockaddr_in addr; struct hostent *server;
server = gethostbyname(str);
if (!server) { /* ... */ };
memcpy(server->h_addr, &addr.sin_addr.s_addr, server->h_length);
```

## getaddrinfo

```
struct addrinfo hints, *result, *rp;
int sockFd = -1;
int ret = getaddrinfo("www.fmph.uniba.sk", "http", &hints, &result);
if (ret != 0) { /*...*/ }
for (rp = result; rp != null; rp = rp->ai_next) {
    sockFd = socket(rp->ai_family, rp->ai_socktype, rp->ai_protocol));
    if (sockFd == -1)
        continue;
    if (connect(sockFd, rp->ai_addr, rp->ai_addrlen) == 0)
        break;
    close(sockFd);
}
freeaddrinfo(result);
if (rp == NULL) { /* could not connect to any of the addresses */ }
else { /* connected... */ }
```

# Copying - problems



<https://www.linuxjournal.com/article/6345>

Image: <https://www.ibm.com/developerworks/library/j-zerocopy/index.html>

# Scatter / gather operations

A final packet consist of multiple parts: ethernet frame header, ip(+tcp) header, user data. Older drivers required continuous buffers, so the parts needed to be copied over to a single buffer.

Newer driver allow scater / gater operations: multiple buffers can be specified for a single packet (either when reading or when writing).

```
struct iovec {  
    void *iov_base;      /* Starting address */  
    size_t iov_len;     /* Number of bytes to transfer */  
};  
  
ssize_t readv(int fd, const struct iovec *iov, int iovcnt);  
ssize_t writev(int fd, const struct iovec *iov, int iovcnt);
```

Note: these still copy between user space and kernel space similar to read and write.

[readv\(2\)](#) [writev\(2\)](#)

# Copying in the kernel

Copy a mmapped file (in\_fd must be "mmap-able"):

```
ssize_t sendfile(int out_fd, int in_fd, off_t *offset, size_t count);
```

Moving data between pipes (or files and pipes):

```
ssize_t splice(int fd_in, loff_t *off_in, int fd_out,
               loff_t *off_out, size_t len, unsigned int flags);
ssize_t tee(int fd_in, int fd_out, size_t len, unsigned int flags);
ssize_t vmsplice(int fd, const struct iovec *iov,
                 unsigned long nr_segs, unsigned int flags);
```

`sendfile(2)` `splice(2)` `vmsplice(2)` `tee(2)`

<https://blogs.oracle.com/linux/post/pipe-and-splice>