

Programovanie v operačných systémoch

08 - Synchronization

Jozef Šiška



Department of Applied Informatics
Comenius University in Bratislava

2025/2026

Problems

Primitives

Memory ordering

Problems 2

Problem

A simple modification of a value in memory ,such as `ref--` is actually three steps on most architectures: fetch, modify and store back:

```
| mov eax,ref      |  
| sub eax, 1      |  
| mov ref, eax    |  
|                   | mov eax,ref  
|                   | sub eax, 1  
|                   | mov ref, eax
```

Non-trivial data structures usually have multiple members, that must be modified "consistently":

```
| data[end] = 42    |  
| end += 1         |  
|                   | data[end] = 47  
|                   | end += 1
```

Problems

- ▶ Execution order, memory (data structure) consistency
- ▶ (Multiple) Read - (Multiple) Write
- ▶ Consumer - Producer
- ▶ Publish Subscribe?

- ▶ Dining Philosophers (ehm)

Primitives

- ▶ Atomic reads, writes
- ▶ TSL, CAS
- ▶ Semaphore
- ▶ Mutex (SpinLock, futext)
- ▶ Wait conditions
- ▶ Monitor
- ▶ Barriers
- ▶ IPC
- ▶ RCU / COW

Spin lock, try 1

```
void lock(int* lock) {  
    while (*lock != 0) ;  
    *lock = 1;  
}  
  
void unlock(int* lock) {  
    *lock = 0;  
}
```

Do or Do not. There is no Try.

Test and set, Compare and swap

Test and set

Write 1 and return the old value "atomically", i.e. only one caller will change it from 0 to 1.

(*tsl* instruction - test and set lock) <https://en.wikipedia.org/wiki/Test-and-set>

Compare and swap

Compare current and a given value, if matches, write a new value. Return the old value if successfull (compare and swap) or a boolean (compare and set).

More general than TSL, can be used to implement semaphores and lock free algorithms.

<https://en.wikipedia.org/wiki/Compare-and-swap>

Fetch and add

Atomically increment (and return). <https://en.wikipedia.org/wiki/Fetch-and-add>

Semaphore

[https://en.wikipedia.org/wiki/Semaphore_\(programming\)](https://en.wikipedia.org/wiki/Semaphore_(programming))

Memory ordering

► Memory ordering

```
a = b = x = y = 0  
x = 1;           y = 1;  
a = y ;         b = x;
```

Memory ordering

► Memory ordering

```
a = b = x = y = 0  
x = 1;           y = 1;  
a = y ;         b = x;
```

► store x, 1 store y, 1 load REG, y load REG, x store a, REG store b, REG

Memory ordering

► Memory ordering

a = b = x = y = 0
x = 1; y = 1;
a = y ; b = x;

- store x, 1 store y, 1
load REG, y load REG, x
store a, REG store b, REG

- load REG, y load REG, x
store x, 1 store y, 1
store a, REG store b, REG

Memory ordering

- ▶ Memory ordering

```
a = b = x = y = 0  
x = 1;           y = 1;  
a = y ;         b = x;
```

- ▶ store x, 1 store y, 1
load REG, y load REG, x
store a, REG store b, REG
- ▶ load REG, y load REG, x
store x, 1 store y, 1
store a, REG store b, REG

- ▶ Sequential consistency

- ▶ Memory barrier

- ▶ Atomic instruction memory semantics

- ▶ <http://preshing.com/20120515/memory-reordering-caught-in-the-act/>
<http://en.cppreference.com/w/cpp/atomic>
http://en.cppreference.com/w/cpp/atomic/memory_order
<https://godbolt.org/z/afWr37j34>

Problems still?

- ▶ deadlock (livelock)
- ▶ priority inversion (priority inheritance)
- ▶ efficiency
- ▶ hard to analyze

Problems still?

- ▶ deadlock (livelock)
- ▶ priority inversion (priority inheritance)
- ▶ efficiency
- ▶ hard to analyze
 - ▶ Mutual exclusion problem
 - ▶ Mutual Exclusion: Only one process/thread can be in the critical section at a time
 - ▶ Progress: No process/thread is forced to wait for an available resource
 - ▶ Bounded Waiting: No process/thread can wait forever for a resource
 - ▶ Lock free, wait free

Other solutions

- ▶ Why is it so hard?
 - ▶ It's how the hardware works...
 - ▶ Current abstractions are still "low" level
 - ▶ We use the wrong paradigms?
- ▶ Higher level apis, managers ("spooler" etc)
 - ▶ tries to hide the details for most things
 - ▶ users most probably need to understand how it works "under the hood" to use it correctly
- ▶ Concentrate on data, not code
 - ▶ Think of what needs to be done with the data / how it moves through the system, not about a sequence of steps that need to be executed
 - ▶ Qt signal / slots
 - ▶ Data flow languages
 - ▶ Immutable data (<https://www.slideshare.net/Kevlin/thinking-outside-the-synchronisation-quadrant/12>)

Other resources

Mutexes, ...

- ▶ memory based, thus mostly used for memory
- ▶ need more work to correctly use between processes

Other resources

- ▶ shared: printer (spooler, print server), hard drives (filesystem), sound card (mixing, pulseaudio), ...
- ▶ harder/not able to share: serial port, most character devices, access to files?
- ▶ data races: creating files and writing to them, creating temporary files
- ▶ file locking (`man flock`), advisory only (processes can still modify files)